

## Approach

The implementation follows a modular design, dividing responsibilities across different classes:

- `FileReaderManager`: Reads words from a file into a list.
- `CharacterFrequencyCounterManager`: Generates a unique frequency-based key for each word.
- `AnagramGrouperManager`: Groups words based on their frequency key.
- `FileWriterManager`: Writes the grouped anagrams back to a file.

## Maintainability

- The modular approach enhances maintainability, making it easy to update or extend specific functionalities independently.
- Methods are well-named and concise, making the code readable and easy to debug.
- Exception handling ensures proper resource management, such as closing files after use.

## Scalability

- The `HashMap` in `AnagramGrouperManager` class ensures efficient lookups and insertions.
- Using a `StringBuilder` in `CharacterFrequencyCounterManager` class avoids excessive string concatenation, improving performance.

## Performance

- The use of a character frequency array (`int[26]`) instead of sorting makes key generation  $O(n)$  rather than  $O(n \log n)$ .
- File operations use buffered reading and writing for better performance.

## External Libraries

No external libraries were used, ensuring simplicity and portability.

## Scalability Considerations

To support large datasets I would:

- Store words and their frequency keys in a database to avoid memory overflows.
- Store files in cloud services to handle large data more efficiently.
- Store words in compressed formats, to reduce I/O overhead.